# One scenario, two prompts

*Appendix · The pilot demo — same scenario for all three vendors, built on their platform, run in our harness*

**PROMPT 1 — THE BUSINESS QUESTION**

*"Across my portfolio, identify products where price increases over the past 52 weeks drove buyer attrition; products that lost households, not just volume. Then build an audience of its lapsed buyers, scale it, and size it for activation."*

The supervisor plans both phases, executes phase one against the Price & Assortment orchestrator (our real orchestrator — returns a top-5 attrition table as an MCP applet), then holds at a human-in-the-loop checkpoint for product selection.

**PROMPT 2 — THE SELECTION AND EXECUTION**

*"Go with the top one — build the audience of its lapsed buyers, scale it, and size it for activation."*

"The top one" resolves from conversation memory. The Liquid Activate orchestrator decomposes into build → scale → size: Build and Size call real MCP tools over private connectivity into Circana APIs; Scale is stubbed. Output of each step feeds the next as structured data, and the flow ends in an activation-ready audience card.

# Target architecture

*Appendix · The pilot demo — platform-native supervisor and orchestrators; solid = real, dashed = stubbed*

*"Across my portfolio, identify products where price increases over the past 52 weeks drove buyer attrition; products that lost households, not just volume. Then build an audience of its lapsed buyers, scale it, and size it for activation."*

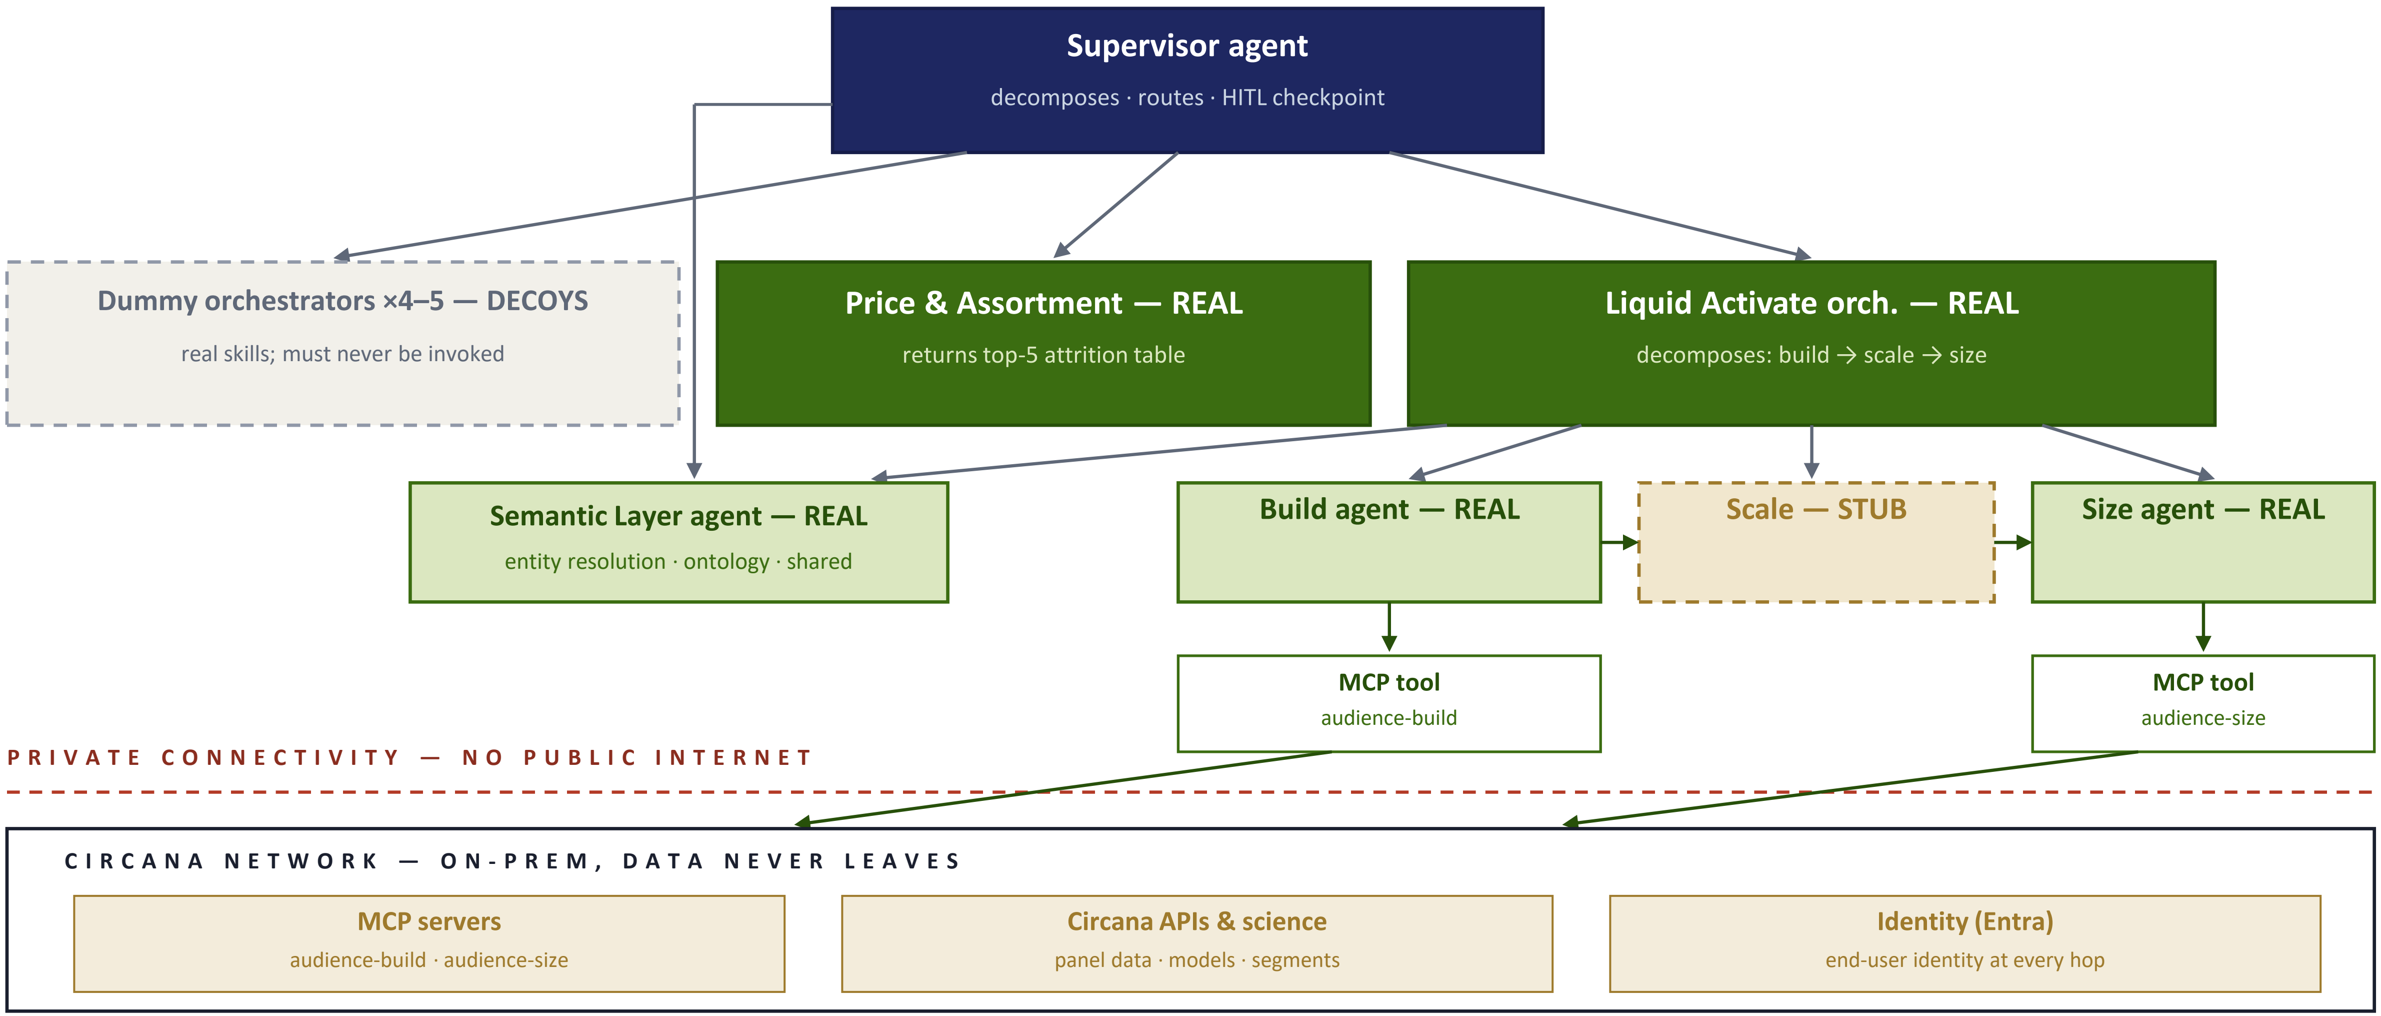

# The two flows

*Appendix · The pilot demo — Flow 1 ends at a checkpoint; Flow 2 runs end to end*

### FLOW 1 — DECOMPOSE, ROUTE, RETURN THE TABLE


*Memory write: five product entities — ids, names, metrics — stored as structured state Flow 2 can resolve “the top one” against.*

### FLOW 2 — BUILD, SCALE, SIZE, END TO END

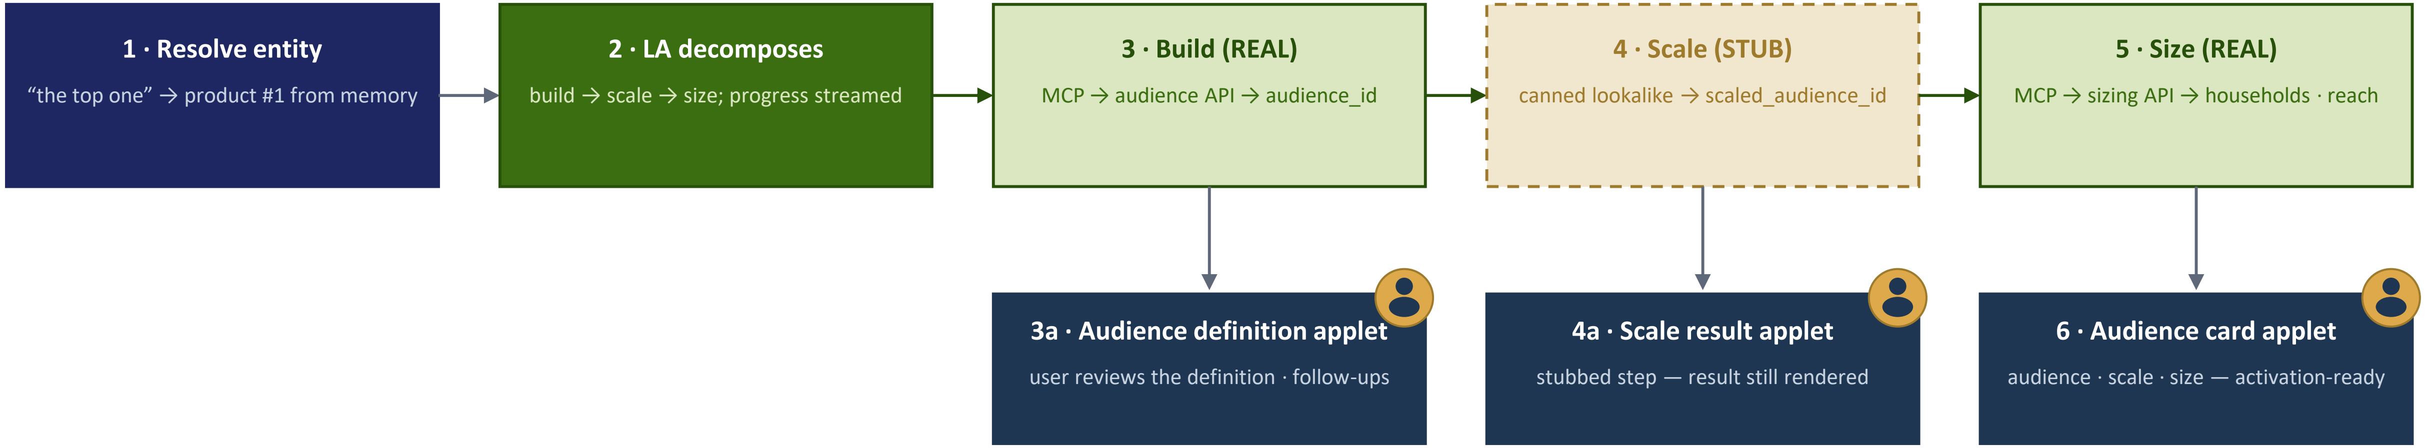


*Gold badge = the user sees or acts in the chat host. Every step’s result reaches the user — stubbed or not — and the test user’s identity rides every hop, verified in our API logs.*

# What Circana provide the hyperscaler

*Appendix · The pilot demo — deliberately minimal: vendors figure out the rest, and how they do is what we score*

| | |
|---|---|
| **1 · The prompt** | Prompt 1, verbatim — nothing else about how to decompose it. Further prompts (including the follow-up selection) are held back and used in scored test runs. |
| **2 · MCP server docs** | Basic documentation for the two Liquid Activate servers we build — audience-build and audience-size. Tool schemas and short descriptions; discovery is part of the test. |
| **3 · Networking & access** | Our side of the private-connectivity setup to the on-prem MCP servers (PrivateLink / Private Endpoint / PSC), with ITO network engineering support. |
| **4 · Chatbot prototype** | The Liquid Activate prototype (or annotated screenshots) to drive UI/UX expectations — the audience-definition MCP app, follow-up buttons, streaming. Their host should meet or beat it. |
| **5 · Decoy skill definitions** | Skills + MCP-server descriptions for 4–5 dummy solution orchestrators, drawn from other Circana solutions. Registered like the real two; any invocation of a decoy = routing failure. |

# Capability coverage — what the demo must showcase

*Appendix · The pilot demo — full sub-topic map in the companion pilot-demo document*

● Core ● Extension ● Discussion

## B1 · Chat host & MCP apps

- ● Applet rendering — table + audience card
- ● Streaming
- ● HITL checkpoint — user selects mid-flow
- ● Error behavior
- ● History & recall
- ● Attachments
- ● Conversation management

## B2 · Prompt orchestration

- ● Multi-part decomposition
- ● Cross-solution spanning
- ● Structured step chaining
- ● Plan transparency
- ● Mid-flow checkpoint
- ● Partial results & failure
- ● Long-running queries
- ● OPS / SME workflow edit

## B3 · Supervisor flow

- ● Native supervisor primitives
- ● Shared state
- ● Routing visibility
- ● Routing precision — decoys never invoked
- ● Guardrail block in-flow
- ● Failure isolation
- ● Per-step cost attribution
- ● Parallelism
- ● LangGraph-over-A2A
- ● Loop control
- ● Verifier agent

## B4 · Memory

- ● Short-term memory — entities across flows
- ● Data residency shown
- ● Memory inspection
- ● Long-term memory — next-day recall
- ● User preferences — e.g. preferred data model
- ● Semantic-layer resolution
- ● Compaction
- ● Consolidation
- ● Retention & deletion

## B5 · MCP private connectivity

- ● Private path verified
- ● Documented network model
- ● Measured latency
- ● Timeout / retry under fault
- ● Spec depth beyond tools
- ● Egress cost lines

## B6 · Identity

- ● User identity at the tool
- ● Token model
- ● Per-user authorization
- ● Delegation boundaries
- ● Entra federation
- ● Audit log
- ● Tenant isolation
- ● Mid-session revocation